

breast_cancer_classification

August 31, 2024

1 BREAST CANCER CLASSIFICATION PROJECT



1.1 Business Problem

Death among women has various factors and one of the leading causes being **Breast Cancer**. A good way of improving the survival rates and giving out personalized treatments depends on whether or not it can be detected early or predicted in advance. The intricacy and diversity of patient data, including age, tumor size, hormone receptor status, and staging, can make it difficult for medical professionals to assess the severity and course of the illness.

A hospital that deals with breast cancer patients want a solution that can be used to classify patients based on their breast cancer characteristics and tell if they are likely to experience a more aggressive cancer. then secondly they would like to be able to predict a patients survival time based oon various factors both medical and demographic. The company can increase the precision of diagnosis and treatment strategies by using machine learning to predict survival months and classify the cancer status (e.g., aggressive vs. non-aggressive). The ultimate goal is to increase survival rates, reduce patient suffering, and optimize the overall cost-effectiveness of breast cancer care.

1.2 Business Understanding

1.2.1 Objectives

1. Create a predictive model to estimate survival time based on medical and demographic factors.

2. Develop a machine learning model to classify patients into categories such as aggressive vs. non-aggressive cancer.

1.2.2 Stakeholders:

1. Patients: Directly affected by the treatment and survival outcomes.
2. Medical Professionals: Need accurate tools to make informed treatment decisions.
3. Hospital Administrators: Focused on improving patient outcomes and optimizing operational costs.

Assumption: The quality and completeness of the data provided are sufficient for building accurate models.

2 Step 1: Data Exploration

```
[1]: import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
import statsmodels.api as sm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, accuracy_score, \
    confusion_matrix, roc_curve, auc, roc_auc_score
from sklearn.tree import DecisionTreeClassifier
from imblearn.over_sampling import SMOTE
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

```
[2]: #Load the dataset
data = pd.read_csv('Breast_Cancer.csv')
data
```

```
[2]:
```

	Age	Race	Marital Status	T Stage	N Stage	6th Stage	\
0	68	White	Married	T1	N1	IIA	
1	50	White	Married	T2	N2	IIIA	
2	58	White	Divorced	T3	N3	IIIC	
3	58	White	Married	T1	N1	IIA	
4	47	White	Married	T2	N1	IIB	
...	...						
4019	62	Other	Married	T1	N1	IIA	
4020	56	White	Divorced	T2	N2	IIIA	
4021	68	White	Married	T2	N1	IIB	
4022	58	Black	Divorced	T2	N1	IIB	
4023	46	White	Married	T2	N1	IIB	

	differentiate	Grade	A Stage	Tumor Size	Estrogen Status	\
0	Poorly differentiated	3	Regional	4	Positive	

1	Moderately differentiated	2	Regional	35	Positive
2	Moderately differentiated	2	Regional	63	Positive
3	Poorly differentiated	3	Regional	18	Positive
4	Poorly differentiated	3	Regional	41	Positive
...
4019	Moderately differentiated	2	Regional	9	Positive
4020	Moderately differentiated	2	Regional	46	Positive
4021	Moderately differentiated	2	Regional	22	Positive
4022	Moderately differentiated	2	Regional	44	Positive
4023	Moderately differentiated	2	Regional	30	Positive

	Progesterone Status	Regional Node Examined	Reginol Node Positive	\
0	Positive	24	1	
1	Positive	14	5	
2	Positive	14	7	
3	Positive	2	1	
4	Positive	3	1	
...
4019	Positive	1	1	
4020	Positive	14	8	
4021	Negative	11	3	
4022	Positive	11	1	
4023	Positive	7	2	

	Survival Months	Status
0	60	Alive
1	62	Alive
2	75	Alive
3	84	Alive
4	50	Alive
...
4019	49	Alive
4020	69	Alive
4021	69	Alive
4022	72	Alive
4023	100	Alive

[4024 rows x 16 columns]

```
[3]: #Look at the basic info of the dataset
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4024 entries, 0 to 4023
Data columns (total 16 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Age                 4024 non-null  int64
```

```

1  Race                4024 non-null  object
2  Marital Status      4024 non-null  object
3  T Stage             4024 non-null  object
4  N Stage             4024 non-null  object
5  6th Stage           4024 non-null  object
6  differentiate       4024 non-null  object
7  Grade               4024 non-null  object
8  A Stage             4024 non-null  object
9  Tumor Size          4024 non-null  int64
10 Estrogen Status     4024 non-null  object
11 Progesterone Status 4024 non-null  object
12 Regional Node Examined 4024 non-null  int64
13 Reginol Node Positive 4024 non-null  int64
14 Survival Months     4024 non-null  int64
15 Status              4024 non-null  object
dtypes: int64(5), object(11)
memory usage: 503.1+ KB

```

```
[4]: #Check for missing values
data.isnull().sum()
```

```

[4]: Age                0
     Race                0
     Marital Status      0
     T Stage             0
     N Stage             0
     6th Stage           0
     differentiate       0
     Grade               0
     A Stage             0
     Tumor Size          0
     Estrogen Status     0
     Progesterone Status 0
     Regional Node Examined 0
     Reginol Node Positive 0
     Survival Months     0
     Status              0
dtype: int64

```

```
[5]: #Check for any duplications in the dataset
data.duplicated().sum()
```

```
[5]: 1
```

```

[6]: #Drop the duplicate values found
data.drop_duplicates(inplace=True)
#check if there still any duplicates left
data.duplicated().sum()

```

```
[6]: 0
```

```
[7]: #Look at the total number of unique values within the dataset
for column in data.columns:
    print(f'{column} has {data[column].nunique()} unique values')
```

```
Age has 40 unique values
Race has 3 unique values
Marital Status has 5 unique values
T Stage has 4 unique values
N Stage has 3 unique values
6th Stage has 5 unique values
differentiate has 4 unique values
Grade has 4 unique values
A Stage has 2 unique values
Tumor Size has 110 unique values
Estrogen Status has 2 unique values
Progesterone Status has 2 unique values
Regional Node Examined has 54 unique values
Reginol Node Positive has 38 unique values
Survival Months has 107 unique values
Status has 2 unique values
```

```
[8]: data.describe()
```

```
[8]:
```

	Age	Tumor Size	Regional Node Examined \
count	4023.000000	4023.000000	4023.000000
mean	53.969923	30.477007	14.358439
std	8.963118	21.121253	8.100241
min	30.000000	1.000000	1.000000
25%	47.000000	16.000000	9.000000
50%	54.000000	25.000000	14.000000
75%	61.000000	38.000000	19.000000
max	69.000000	140.000000	61.000000

	Reginol Node Positive	Survival Months
count	4023.000000	4023.000000
mean	4.158837	71.301765
std	5.109724	22.923009
min	1.000000	1.000000
25%	1.000000	56.000000
50%	2.000000	73.000000
75%	5.000000	90.000000
max	46.000000	107.000000

3 Step 2: Exploratory Data Analysis

```
[9]: #here we are going to group the data to help us later on when performing
      ↪ regression and classification
numerical_cols = ['Age', 'Tumor Size', 'Regional Node Examined', 'Reginol Node_
      ↪ Positive', 'Survival Months']
categorical_cols = ['Race', 'Marital Status', 'T Stage ', 'N Stage', '6th_
      ↪ Stage', 'differentiate', 'Grade', 'A Stage', 'Estrogen Status',_
      ↪ 'Progesterone Status', 'Status']
```

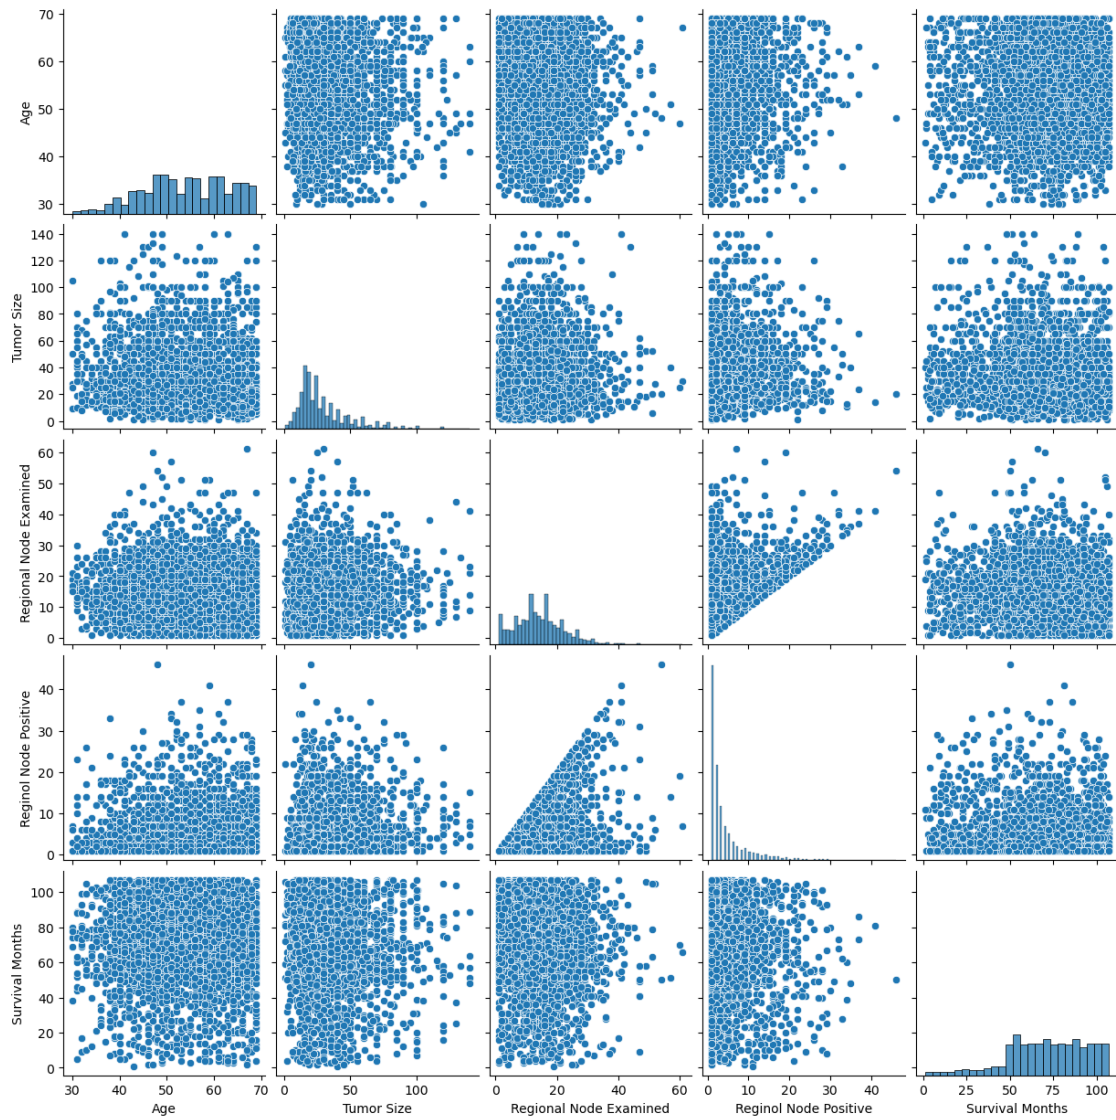
3.1 Linear Regression


for this part we want to create a simple model that will see the numerical data being fitted into a linear regression in order to predict the number of survival months

```
[10]: #Use the numerical_cols as the dataframe to use for creating the Linear_
      ↪ regression
data_numerals = data[numerical_cols]
#create a pairplot for the dataframe
sns.pairplot(data_numerals)
plt.suptitle('Pair Plot of data_numerals Columns', y=1.02)
plt.show()
# Correlation Matrix
plt.figure(figsize=(10, 6))
sns.heatmap(data_numerals.corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix')
plt.show()
```

```
c:\Users\User\.conda\envs\learn-env\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\User\.conda\envs\learn-env\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\User\.conda\envs\learn-env\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\User\.conda\envs\learn-env\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\User\.conda\envs\learn-env\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
```

Pair Plot of data_numerals Columns





```
[11]: #Initialize the scaler
scaler = StandardScaler()

#Select the feature and target variables
#define the target
y_lin = data_numerals['Survival Months']

#fit and transform the features
X_lin = scaler.fit_transform(data_numerals.drop('Survival Months', axis = 1))

[12]: #Now let's create a linear regression model that will cater for the numerical
      ↪ data that we have in the data set
lg_model = sm.OLS(y_lin, sm.add_constant(X_lin))
lg_results = lg_model.fit()
lg_results.summary()
```

[12]:

Dep. Variable:	Survival Months	R-squared:	0.023
Model:	OLS	Adj. R-squared:	0.022
Method:	Least Squares	F-statistic:	23.57
Date:	Sat, 31 Aug 2024	Prob (F-statistic):	2.73e-19
Time:	04:04:39	Log-Likelihood:	-18262.
No. Observations:	4023	AIC:	3.653e+04
Df Residuals:	4018	BIC:	3.657e+04
Df Model:	4		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	71.3018	0.357	199.491	0.000	70.601	72.003
x1	-0.2444	0.359	-0.681	0.496	-0.948	0.459
x2	-1.3458	0.370	-3.641	0.000	-2.071	-0.621
x3	0.9211	0.393	2.347	0.019	0.152	1.691
x4	-3.1517	0.402	-7.832	0.000	-3.941	-2.363

Omnibus:	169.896	Durbin-Watson:	1.987
Prob(Omnibus):	0.000	Jarque-Bera (JB):	191.623
Skew:	-0.535	Prob(JB):	2.45e-42
Kurtosis:	2.980	Cond. No.	1.66

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

1. Model Summary Dep. Variable (Dependent Variable): Survival Months

The variable that the model is trying to predict. R-squared: 0.023

R-squared is the proportion of the variance in the dependent variable that is predictable from the independent variables. An R-squared of 0.023 means that only 2.3% of the variability in Survival Months is explained by the model. This is a very low value, indicating that the model doesn't explain much of the variability in survival months. Adj. R-squared: 0.022

The adjusted R-squared accounts for the number of predictors in the model and adjusts for the potential overfitting. It's slightly lower than the R-squared, indicating that adding more predictors didn't improve the model much. F-statistic: 23.57

The F-statistic tests whether at least one of the predictors is statistically significant. A higher F-statistic indicates that the model is a better fit than a model with no predictors. Prob (F-statistic): 2.73e-19

This is the p-value associated with the F-statistic. A very small p-value (close to zero) indicates that the model as a whole is statistically significant.

2. Coefficients and Statistical Significance

Const (Intercept): 71.3018

This is the expected value of Survival Months when all independent variables are zero. The constant term is highly significant with a p-value of 0.000.

x1, x2, x3, x4 (Independent Variables):

x1 Coefficient: -0.2444

The relationship between x1 and Survival Months is negative, but with a p-value of 0.496, it is not statistically significant. This suggests that x1 does not have a meaningful impact on Survival Months.

x2 Coefficient: -1.3458

The relationship between x2 and Survival Months is negative and statistically significant (p-value = 0.000). For every unit increase in x2, Survival Months decreases by approximately 1.35 months.

x3 Coefficient: 0.9211

The relationship between x3 and Survival Months is positive and statistically significant (p-value = 0.019). This suggests that as x3 increases, Survival Months increases by approximately 0.92 months.

x4 Coefficient: -3.1517

The relationship between x4 and Survival Months is strongly negative and highly significant (p-value = 0.000). For every unit increase in x4, Survival Months decreases by approximately 3.15 months.

Standard Error (std err):

Reflects the accuracy of the coefficient estimates. Smaller standard errors indicate more precise estimates.

t-statistic (t) and $P > |t|$ (p-value):

The t-statistic tests whether a coefficient is significantly different from zero. The p-value tells you the probability that the coefficient is actually zero (i.e., not significant).

Coefficients with p-values less than 0.05 are generally considered statistically significant.

The condition number tests for multicollinearity (high correlation between independent variables). A value below 30 suggests no serious multicollinearity issues.

Conclusion

Model Fit: The R-squared and adjusted R-squared values are low, indicating that the model does not explain much of the variability in Survival Months.

Significance: While the model overall is statistically significant (based on the F-statistic), not all variables are contributing meaningfully (e.g., x1 is not significant).

Coefficients: Some variables (x2, x3, x4) have significant coefficients, suggesting they do impact survival months, but the effect sizes vary.

Assumptions: The normality tests suggest that the residuals may not be normally distributed, which could affect the validity of the model's inferences.

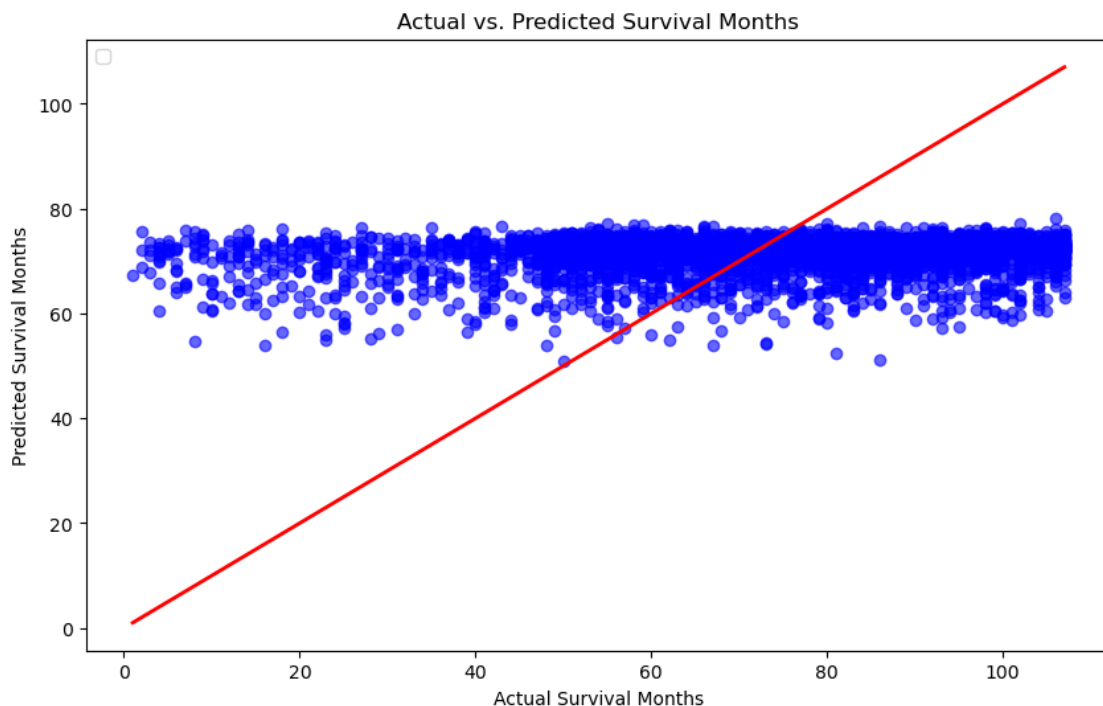
Given these results, the model might require refinement, such as exploring non-linear relationships, interaction effects, or considering other variables that might better explain the variation in survival months.

```
[13]: # Predicted values
y_pred = lg_results.predict(sm.add_constant(X_lin))

plt.figure(figsize=(10, 6))

plt.scatter(y_lin, y_pred, color='blue', alpha=0.6)
plt.plot([y_lin.min(), y_lin.max()], [y_lin.min(), y_lin.max()], color='red',
        linewidth=2)
plt.xlabel('Actual Survival Months')
plt.ylabel('Predicted Survival Months')
plt.title('Actual vs. Predicted Survival Months')
plt.legend(loc='upper left')
plt.show()
```

No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.



Interpretation:

Points on the Red Line: Data points on this line indicate perfect predictions.

Points Above the Red Line: Data points above the line indicate overestimations by the model (the predicted survival is higher than the actual survival).

Points Below the Red Line: Data points below the line indicate underestimations by the model (the predicted survival is lower than the actual survival).

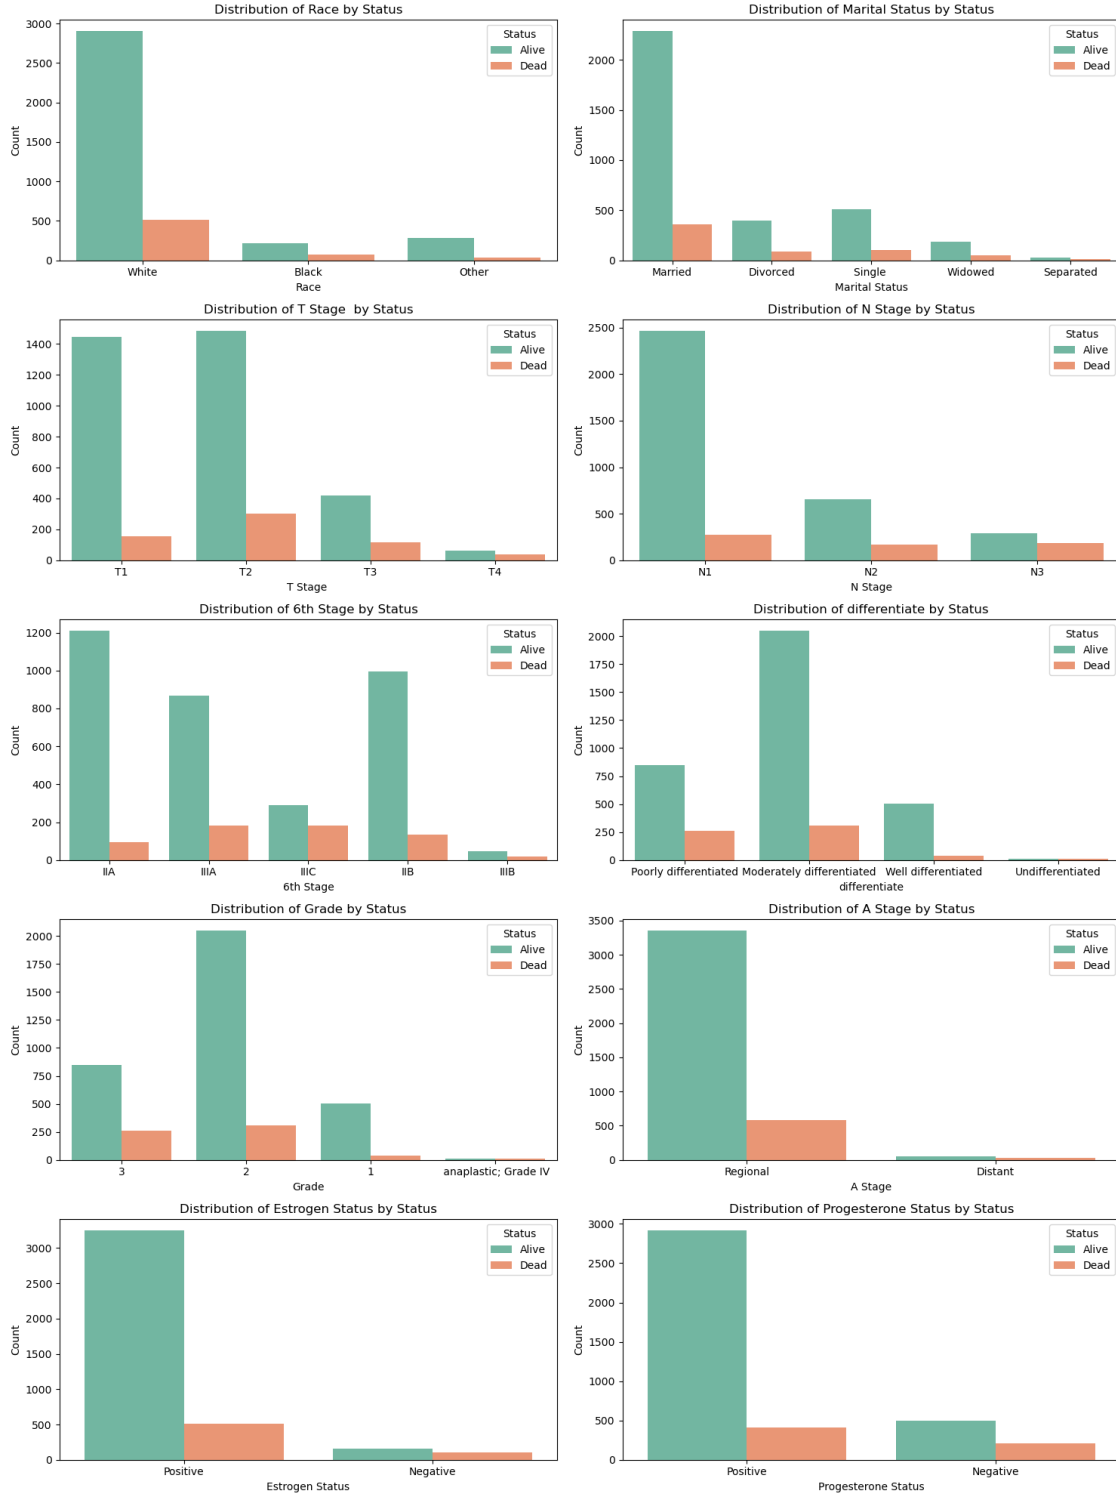
Cluster of Blue Dots: Most of the data points are clustered horizontally between 20 and 80 on the x-axis, showing a concentration of actual survival months in that range. However, these points spread vertically, indicating that the model's predictions vary significantly from the actual survival months in this range.

Overall Insight:

The graph shows that while the model predicts survival months, the predictions deviate quite a bit from the actual values, as seen by the spread of points around the red line. The red line serves as a reference for ideal predictions, but most of the points do not lie on this line, suggesting that the model may not be very accurate in predicting the exact survival months.

4 CLASSIFICATION MODELS

```
[14]: #here we create a plot of the distribution of the categorical columns with
      ↪repsect to status
columns_to_plot = data[categorical_cols].drop(columns='Status',axis=1)
plt.figure(figsize=(15, 20))
for i, col in enumerate(columns_to_plot, 1):
    plt.subplot(5, 2, i) # Adjust based on the number of plots (5 rows, 2
    ↪columns here)
    sns.countplot(x=col, hue='Status', data=data[categorical_cols],
    ↪palette='Set2')
    plt.title(f'Distribution of {col} by Status')
    plt.xlabel(col)
    plt.ylabel('Count')
plt.tight_layout()
plt.show()
```



```
[15]: #Define the data we are using for classification i.e. the categorical_cols
X = data[categorical_cols]
```

```
#One-Hot encode the variables
X=pd.get_dummies(X, drop_first=True, dtype=int)
X.head(25)
```

```
[15]:
```

	Race_Other	Race_White	Marital Status_Married	Marital Status_Separated	\
0	0	1	1	0	
1	0	1	1	0	
2	0	1	0	0	
3	0	1	1	0	
4	0	1	1	0	
5	0	1	0	0	
6	0	1	1	0	
7	0	1	1	0	
8	0	1	0	0	
9	0	1	1	0	
10	0	1	0	0	
11	0	1	1	0	
12	0	1	1	0	
13	0	1	1	0	
14	0	1	0	0	
15	0	1	1	0	
16	0	1	0	0	
17	0	1	1	0	
18	0	0	0	0	
19	0	1	0	0	
20	1	0	1	0	
21	0	1	1	0	
22	0	1	0	0	
23	0	1	1	0	
24	0	1	1	0	

	Marital Status_Single	Marital Status_Widowed	T Stage _T2	T Stage _T3	\
0	0	0	0	0	
1	0	0	1	0	
2	0	0	0	1	
3	0	0	0	0	
4	0	0	1	0	
5	1	0	0	0	
6	0	0	0	0	
7	0	0	1	0	
8	0	0	0	0	
9	0	0	0	0	
10	0	1	0	0	
11	0	0	0	1	
12	0	0	1	0	
13	0	0	0	0	
14	0	0	1	0	

15	0	0	0	0
16	1	0	1	0
17	0	0	1	0
18	0	0	1	0
19	0	0	1	0
20	0	0	1	0
21	0	0	1	0
22	1	0	1	0
23	0	0	0	1
24	0	0	0	0

	T Stage _T4	N Stage_N2	...	differentiate_Poorly	differentiated \
0	0	0	...		1
1	0	1	...		0
2	0	0	...		0
3	0	0	...		1
4	0	0	...		1
5	0	0	...		0
6	0	0	...		0
7	0	0	...		0
8	1	0	...		1
9	1	0	...		0
10	0	0	...		0
11	0	0	...		1
12	0	0	...		1
13	0	1	...		1
14	0	0	...		0
15	0	0	...		0
16	0	0	...		0
17	0	0	...		0
18	0	0	...		0
19	0	0	...		0
20	0	0	...		0
21	0	1	...		0
22	0	0	...		1
23	0	0	...		1
24	0	0	...		1

	differentiate_Undifferentiated	differentiate_Well	differentiated \
0		0	0
1		0	0
2		0	0
3		0	0
4		0	0
5		0	0
6		0	1
7		0	0

8	0	0
9	0	1
10	0	0
11	0	0
12	0	0
13	0	0
14	0	0
15	0	0
16	0	0
17	0	0
18	0	0
19	0	0
20	0	1
21	0	0
22	0	0
23	0	0
24	0	0

	Grade_1	Grade_2	Grade_3	A	Stage_Regional	Estrogen	Status_Positive	\
0	0	0	1		1		1	
1	0	1	0		1		1	
2	0	1	0		1		1	
3	0	0	1		1		1	
4	0	0	1		1		1	
5	0	1	0		1		1	
6	1	0	0		1		1	
7	0	1	0		1		1	
8	0	0	1		1		1	
9	1	0	0		0		1	
10	0	1	0		1		1	
11	0	0	1		1		0	
12	0	0	1		1		1	
13	0	0	1		1		1	
14	0	1	0		1		1	
15	0	1	0		1		1	
16	0	1	0		1		1	
17	0	1	0		1		1	
18	0	1	0		1		1	
19	0	1	0		1		0	
20	1	0	0		1		1	
21	0	1	0		1		1	
22	0	0	1		1		1	
23	0	0	1		1		1	
24	0	0	1		1		1	

	Progesterone	Status_Positive	Status_Dead
0		1	0

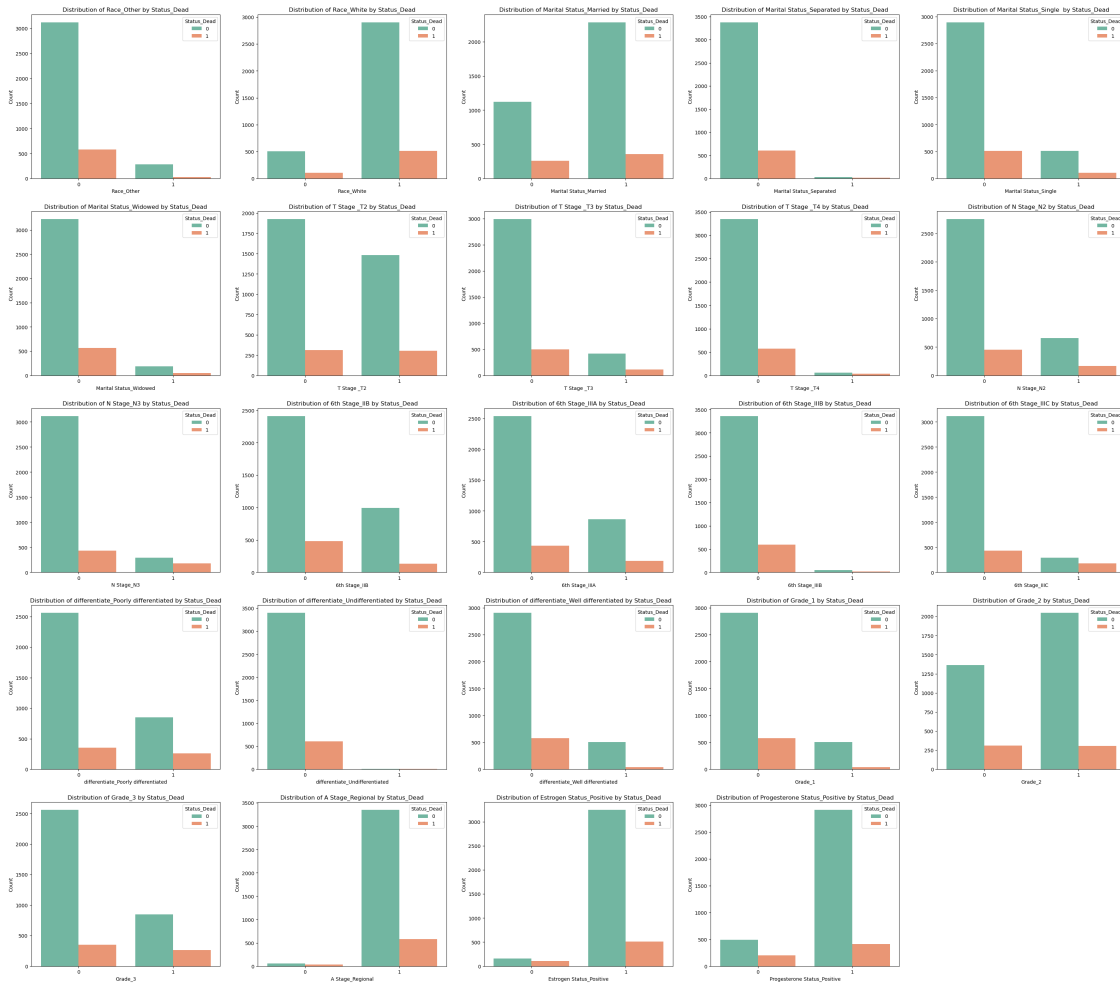
1	1	0
2	1	0
3	1	0
4	1	0
5	1	0
6	1	0
7	1	1
8	1	0
9	1	0
10	1	1
11	0	0
12	1	0
13	1	0
14	1	0
15	1	0
16	1	0
17	1	0
18	1	0
19	0	0
20	1	0
21	1	0
22	1	0
23	0	0
24	1	0

[25 rows x 25 columns]

```
[16]: # Visualize the encoded variables above

plt.figure(figsize=(40, 35))
for i, col in enumerate(X.columns[:-1], 1): # Exclude 'Status_Dead'
    plt.subplot(5, 5, i) # Adjust based on the number of plots (5 rows, 5
    ↪ columns here)
    sns.countplot(x=col, hue='Status_Dead', data=X, palette='Set2')
    plt.title(f'Distribution of {col} by Status_Dead')
    plt.xlabel(col)
    plt.ylabel('Count')

plt.show()
```



```
[17]: #select the features(X_cat) and the target(y_cat)
y_cat = X['Status_Dead']
X_cat = X.drop(columns='Status_Dead',axis = 1)
```

```
[18]: #create the test and train sets
X_train, X_test, y_train, y_test = train_test_split(X_cat, y_cat, test_size=0.
↪3, random_state=1)
```

4.1 The Logistic Regression Model

```
[19]: #create the logistic regression model
logreg = LogisticRegression()
logreg.fit(X_train, y_train)
```

```
[19]: LogisticRegression()
```

```

[20]: #make predictions for the model
y_pred_train = logreg.predict(X_train)
y_pred_test = logreg.predict(X_test)

[21]: #Model Evaluation
#Creating a classification report
train_class_report = classification_report(y_train, y_pred_train)
test_class_report = classification_report(y_test, y_pred_test)
#Display the results of the classification report
print('The outcome of the training classification report is:')
print(train_class_report)
print('\n\n')
print('The outcome of the test classification report is:')
print(test_class_report)
print('\n\n')

#Let's create a confusion matrix that can help provide a summary of prediction
↪ results
train_conf_mat = confusion_matrix(y_train, y_pred_train)
test_conf_mat = confusion_matrix(y_test, y_pred_test)
#Display the confusion matrix results for the train and test sets
print('The confusion matrix for the train set is:')
print(train_conf_mat)
print('\n\n')
print('The confusion matrix for the test set is:')
print(test_conf_mat)
print('\n\n')

#Now let's look at the accuracy score for both sets
train_acc = accuracy_score(y_train, y_pred_train)
test_acc = accuracy_score(y_test, y_pred_test)
#Display the results of the accuracy score
print('The accuracy score for the train set is:', train_acc)
print('The accuracy score for the test set is:', test_acc)

#Finally we should consider the ROC curve and the AUC
#start by getting probability estimates of the positive class
y_train_prob = logreg.predict_proba(X_train)[:, 1]
y_test_prob = logreg.predict_proba(X_test)[:, 1]
#ROC curve for the training data
fpr_train, tpr_train, threshold_train = roc_curve(y_train, y_train_prob)
roc_auc_train = auc(fpr_train, tpr_train)
#ROC curve for the test set
fpr_test, tpr_test, threshold_test = roc_curve(y_test, y_test_prob)
roc_auc_test = auc(fpr_test, tpr_test)
#plot the ROC curve
plt.figure()

```

```
plt.plot(fpr_train, tpr_train, color='blue', lw=2, label=f'Training ROC curve_
↳(area = {roc_auc_train:.2f})')
plt.plot(fpr_test, tpr_test, color='red', lw=2, label=f'Test ROC curve (area =_
↳{roc_auc_test:.2f})')
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Logistic Regression')
plt.legend(loc="lower right")
plt.show()
```

The outcome of the training classification report is:

	precision	recall	f1-score	support
0	0.86	0.99	0.92	2390
1	0.59	0.10	0.17	426
accuracy			0.85	2816
macro avg	0.73	0.54	0.54	2816
weighted avg	0.82	0.85	0.81	2816

The outcome of the test classification report is:

	precision	recall	f1-score	support
0	0.86	0.99	0.92	1017
1	0.70	0.14	0.23	190
accuracy			0.86	1207
macro avg	0.78	0.56	0.57	1207
weighted avg	0.84	0.86	0.81	1207

The confusion matrix for the train set is:

```
[[2362  28]
 [ 385  41]]
```


The confusion matrix for the test set is:

```
[[1006  11]
```

[164 26]]

The accuracy score for the train set is: 0.8533380681818182

The accuracy score for the test set is: 0.8550124275062138



For the training set:

The model has high precision (0.86) for class 0, meaning it correctly predicted most instances labeled as 0. The recall for class 1 is low (0.10), indicating that the model missed many actual positive instances. The overall accuracy is 0.85, which is reasonably good. For the testing set:

The precision and recall for class 0 have improved slightly. The precision for class 1 has also improved, but the recall remains low. The overall accuracy has increased slightly to 0.86. Confusion Matrices:

A confusion matrix shows the number of instances predicted for each class compared to their actual classes. For the training set:

Many instances of class 0 were correctly predicted (2362), but there were 28 false positives. Many instances of class 1 were incorrectly predicted as class 0 (385), and only 41 were correctly predicted. For the testing set:

Similar patterns are observed, with more correct predictions for class 0 and fewer correct predictions for class 1. Accuracy Scores:

The accuracy score measures the overall proportion of correct predictions. Overall:

The model is performing reasonably well on both training and testing sets, with higher accuracy for the testing set. However, the model struggles to correctly predict instances of class 1, especially in the training set. Further analysis and potential improvements might be necessary to address the class imbalance and improve the model's performance on class 1.

Graph Interpretation:

Shape of the curve: A ROC curve ideally should be as close to the top-left corner as possible. This indicates that the model has high TPR and low FPR, which is desirable.

Comparison of curves: In this graph, the test ROC curve is slightly above the training ROC curve, suggesting that the model might be slightly underfitting the training data. However, both curves are significantly above the random classifier line, indicating that the model is performing better than random guessing.

Area Under the Curve (AUC): The AUC is a metric that quantifies the overall performance of a model. A higher AUC indicates better performance. In this case, the AUC for both training and testing curves is above 0.5, indicating that the model is performing better than random guessing. The test AUC is slightly higher than the training AUC, which is a good sign.

Overall, the ROC curve suggests that the logistic regression model is performing reasonably well on both training and testing data, and is significantly better than a random classifier. The slight underfitting indicated by the difference between the training and testing curves might be addressed by adjusting the model's hyperparameters or collecting more data.

4.2 Decision tree model

4.2.1 Vanilla Model of the Decision Tree

```
[22]: #Initiate the classifier
      clf = DecisionTreeClassifier(random_state=42)

      #fit it on the training data
      clf.fit(X_train, y_train)
```

```
[22]: DecisionTreeClassifier(random_state=42)
```

```
[23]: # Make predictions
      y_tr_pred = clf.predict(X_train)
      y_te_pred = clf.predict(X_test)
```

```
[24]: # Evaluate the model
      #Look at the classification report
      print('The classification report for the training set is:')
      print(classification_report(y_train, y_tr_pred))
      print('\n')
      print('The classification report for the test set is:')
```

```

print(classification_report(y_test, y_te_pred))
print('\n\n')
#Show accuracy for both sets
print('the accuracy for the training set is:', accuracy_score(y_train,
    ↪y_tr_pred))
print('the accuracy for the test set is:', accuracy_score(y_test, y_te_pred))
print('\n')
#Create the ROC curve and AUC for both sets
#start by getting the probability for each of the sets
clf_train_prob = clf.predict_proba(X_train)[: , 1]
clf_test_prob = clf.predict_proba(X_test)[: , 1]
#calculate the ROC and auc
train_fpr, train_tpr, train_threshold = roc_curve(y_train, clf_train_prob)
train_auc = auc(train_fpr, train_tpr)
test_fpr, test_tpr, test_threshold = roc_curve(y_test, clf_test_prob)
test_auc = auc(test_fpr, test_tpr)
#plot the ROC curve
plt.figure()
plt.plot(train_fpr, train_tpr, color='blue', label=f'Train AUC = {train_auc:.
    ↪2f}')
plt.plot(test_fpr, test_tpr, color='red', linestyle='--', label=f'Test AUC =
    ↪{test_auc:.2f}')
plt.plot([0, 1], [0, 1], color='grey', linestyle=':')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()

```

The classification report for the training set is:

	precision	recall	f1-score	support
0	0.88	0.99	0.94	2390
1	0.89	0.26	0.40	426
accuracy			0.88	2816
macro avg	0.89	0.63	0.67	2816
weighted avg	0.88	0.88	0.85	2816

The classification report for the test set is:


	precision	recall	f1-score	support
0	0.87	0.97	0.91	1017
1	0.53	0.21	0.30	190

accuracy			0.85	1207
macro avg	0.70	0.59	0.60	1207
weighted avg	0.81	0.85	0.82	1207

\n

the accuracy for the training set is: 0.8831676136363636

th eaccuracy for the test set is: 0.8458989229494615



Insights from the results 1. Precision, Recall and F1_score class 0 (Majority Class):

Training Set: Precision = 0.88, Recall = 0.99, F1-Score = 0.94

Test Set: Precision = 0.87, Recall = 0.97, F1-Score = 0.91

The model performs very well for class 0, with high precision, recall, and F1-score in both the training and test sets.

Class 1 (Minority Class):

Training Set: Precision = 0.89, Recall = 0.26, F1-Score = 0.40

Test Set: Precision = 0.53, Recall = 0.21, F1-Score = 0.30

The performance for class 1 is notably weaker, particularly in terms of recall and F1-score. The

2. Imbalance in Performance

The disparity between the performance on class 0 and class 1 suggests that the model is biased towards predicting the majority class. This is common in imbalanced datasets where one class significantly outnumbers the other. The model's high accuracy is driven by its performance on class 0, but its ability to detect class 1 is poor.

3. Conclusion

While the model performs well overall, its ability to detect the minority class (class 1) is lacking. There is a need to address the class imbalance in the data.

4.2.2 Addressing class imbalance

There are several ways in which the class imbalance can be addressed and the one of the methods we are going to consider first is the use of SMOTE. This is a resampling technique where we use oversampling for the minority class.

```
[25]: #Apply SMOTE to the training data
smote = SMOTE(random_state=42)
smote_X_train, smote_y_train = smote.fit_resample(X_train,y_train)
```

```
[26]: #Train a model
clf.fit(smote_X_train, smote_y_train)
```

```
[26]: DecisionTreeClassifier(random_state=42)
```

```
[27]: #create predictions for train and test sets
smote_train_pred = clf.predict(smote_X_train)
smote_test_pred = clf.predict(X_test)
```

```
[28]: print("Classification Report for Training Set:")
print(classification_report(smote_y_train, smote_train_pred))
print("Classification Report for Test Set:")
print(classification_report(y_test, smote_test_pred))
print('\n\n')
#Show accuracy for both sets
print('the accuracy for the training set is:',accuracy_score(smote_y_train,
    ↪smote_train_pred))
print('the accuracy for the test set is:', accuracy_score(y_test,
    ↪smote_test_pred))
print('\n')
#Create the ROC curve and AUC for both sets
#start by getting the probability for each of the sets
smote_train_prob = clf.predict_proba(smote_X_train)[: , 1]
smote_test_prob = clf.predict_proba(X_test)[: , 1]
#calculate the ROC and auc
```

```

smote_train_fpr, smote_train_tpr, smote_train_threshold =  

    ↪roc_curve(smote_y_train, smote_train_prob)  

train_auc = auc(smote_train_fpr, smote_train_tpr)  

smote_test_fpr, smote_test_tpr, smote_test_threshold = roc_curve(y_test,  

    ↪smote_test_prob)  

test_auc = auc(smote_test_fpr, smote_test_tpr)  

#plot the ROC roc_curve  

plt.figure()  

plt.plot(smote_train_fpr, smote_train_tpr, color='blue', label=f'Train AUC =  

    ↪{train_auc:.2f}')  

plt.plot(smote_test_fpr, smote_test_tpr, color='red', linestyle='--',  

    ↪label=f'Test AUC = {test_auc:.2f}')  

plt.plot([0, 1], [0, 1], color='grey', linestyle=':')  

plt.xlabel('False Positive Rate')  

plt.ylabel('True Positive Rate')  

plt.title('ROC Curve')  

plt.legend(loc='lower right')  

plt.show()

```

Classification Report for Training Set:


	precision	recall	f1-score	support
0	0.76	0.79	0.78	2390
1	0.78	0.75	0.77	2390
accuracy			0.77	4780
macro avg	0.77	0.77	0.77	4780
weighted avg	0.77	0.77	0.77	4780

Classification Report for Test Set:

	precision	recall	f1-score	support
0	0.89	0.74	0.81	1017
1	0.27	0.49	0.35	190
accuracy			0.71	1207
macro avg	0.58	0.62	0.58	1207
weighted avg	0.79	0.71	0.74	1207

the accuracy for the training set is: 0.7713389121338912

th eaccuracy for the test set is: 0.7050538525269263



Now while the use of SMOTE has addressed the issue of class imbalance , the model may still need further refinement to generalize better on unseen data. Here below are some key highlights supporting this:

1. ROC Curve Analysis:

Train AUC: 0.86 Test AUC: 0.62 The train AUC has increased slightly from 0.83 to 0.86, indicating that the model is better at distinguishing between classes on the training data after SMOTE was applied. The test AUC remains the same at 0.62, suggesting that the model's ability to generalize to unseen data hasn't improved significantly.

2. Classification Report Summary:

Training Set (After SMOTE): Accuracy: 0.77 Precision for class 1: 0.78 Recall for class 1: 0.75 F1-score for class 1: 0.77

Test Set (After SMOTE): Accuracy: 0.71 Precision for class 1: 0.27 Recall for class 1: 0.49 F1-score for class 1: 0.35 Interpretation:

3. Impact of SMOTE:

SMOTE has balanced the classes in the training data, leading to a more balanced precision and recall for both classes in the training set. For the test set, the recall for class 1 (the minority class) improved from 0.21 to 0.49, indicating the model's improved ability to identify positive

cases. However, the precision for class 1 decreased from 0.53 to 0.27, showing that the model is now making more false positive errors.

4. Accuracy Decrease:

The overall accuracy on both training and test sets has decreased compared to before applying SMOTE (from 0.88 to 0.77 on the training set and from 0.85 to 0.71 on the test set). This is typical when using SMOTE as the model is forced to learn more about the minority class, which can reduce accuracy but improve the balance of errors between classes. Model Generalization:

The test AUC remains unchanged at 0.62, which suggests that while the model's ability to handle class imbalance has improved, its overall discriminatory power on unseen data hasn't.

Now we want to try and tune this model by trying to fine tune the hyperparameters and then have a final model.

4.2.3 Hyperparameter tuning I

We want to start by looking for the maximum depth of the model

```
[29]: # Initialize lists to store results for different depths
depths = range(1, 21)
train_aucs = []
test_aucs = []
train_accuracies = []
test_accuracies = []

# Loop over different values of max_depth
for max_depth in depths:
    # Calculate ROC AUC scores for both sets
    train_auc = roc_auc_score(smote_y_train, smote_train_prob)
    test_auc = roc_auc_score(y_test, smote_test_prob)

    # Calculate accuracy for both sets
    train_accuracy = clf.score(smote_X_train, smote_y_train)
    test_accuracy = clf.score(X_test, y_test)

    # Store the results
    train_aucs.append(train_auc)
    test_aucs.append(test_auc)
    train_accuracies.append(train_accuracy)
    test_accuracies.append(test_accuracy)

# Step 5: Plot AUC scores for different max_depth values
plt.figure(figsize=(10, 6))
plt.plot(depths, train_aucs, marker='o', label='Train AUC')
plt.plot(depths, test_aucs, marker='o', label='Test AUC')
plt.xlabel('max_depth')
plt.ylabel('AUC Score')
```

```

plt.title('AUC Score vs. max_depth for Decision Tree')
plt.legend()
plt.grid(True)
plt.show()

# Step 6: Plot Accuracy for different max_depth values
plt.figure(figsize=(10, 6))
plt.plot(depths, train_accuracies, marker='o', label='Train Accuracy')
plt.plot(depths, test_accuracies, marker='o', label='Test Accuracy')
plt.xlabel('max_depth')
plt.ylabel('Accuracy')
plt.title('Accuracy vs. max_depth for Decision Tree')
plt.legend()
plt.grid(True)
plt.show()

# Step 7: Output classification report for the best max_depth
best_depth = depths[test_aucs.index(max(test_aucs))]
print(f"Best max_depth based on Test AUC: {best_depth}")

best_clf = DecisionTreeClassifier(max_depth=best_depth, random_state=42)
best_clf.fit(smote_X_train, smote_y_train)

smote_train_pred = best_clf.predict(smote_X_train)
smote_test_pred = best_clf.predict(X_test)

print("Classification Report for Training Set:")
print(classification_report(smote_y_train, smote_train_pred))

print("Classification Report for Test Set:")
print(classification_report(y_test, smote_test_pred))

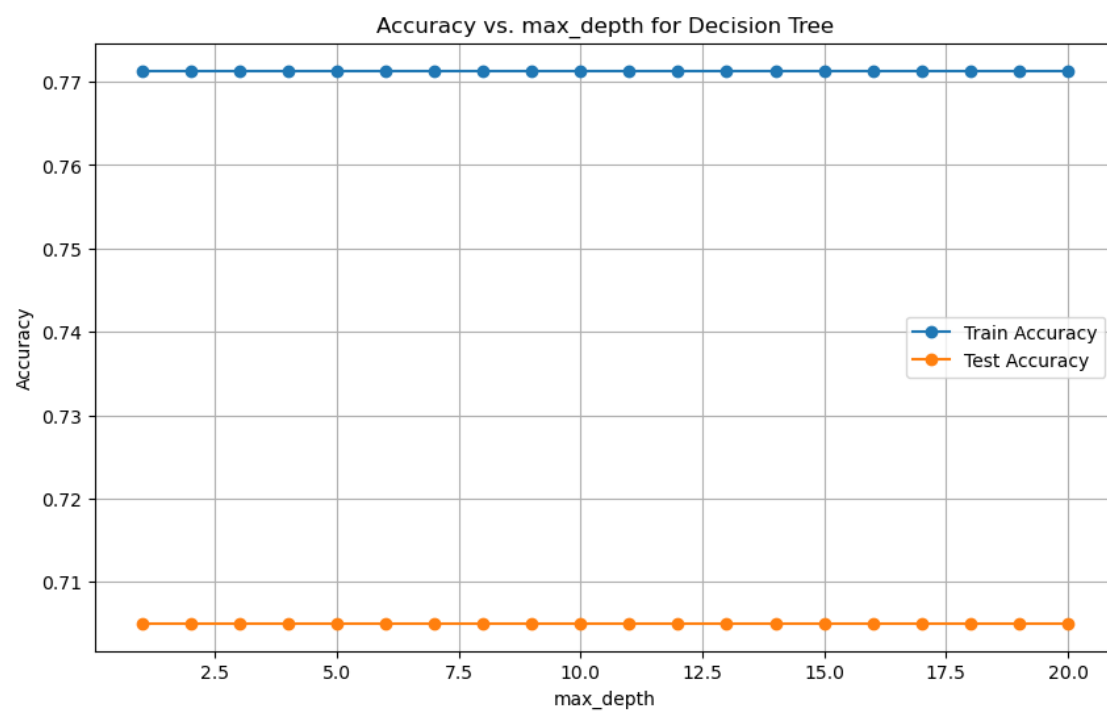
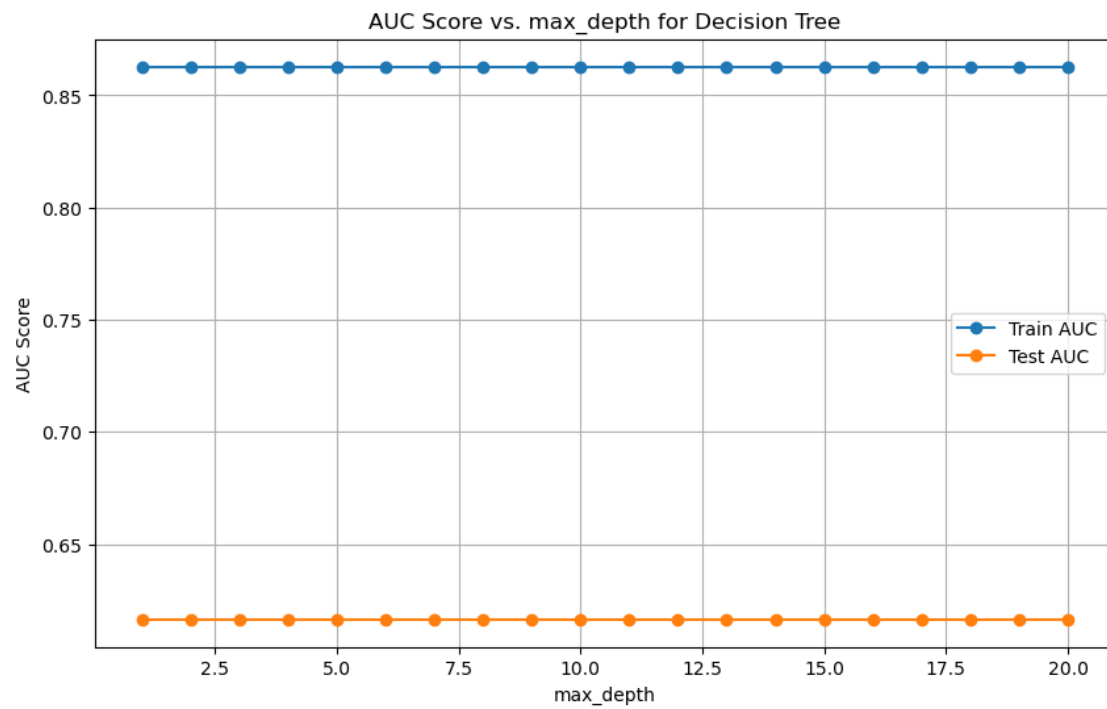
train_auc = roc_auc_score(smote_y_train, best_clf.predict_proba(smote_X_train)[:
    ↪, 1])
test_auc = roc_auc_score(y_test, best_clf.predict_proba(X_test)[: , 1])

print(f"Train AUC: {train_auc:.2f}")
print(f"Test AUC: {test_auc:.2f}")

# Plot ROC curves
plt.figure(figsize=(10, 6))
plt.plot(fpr_train, tpr_train, color='blue', label='Train ROC Curve')
plt.plot(fpr_test, tpr_test, color='green', label='Test ROC Curve')
plt.plot([0, 1], [0, 1], color='gray', linestyle='--', label='Random_
    ↪Classifier')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')

```

```
plt.title('ROC Curve for Decision Tree')
plt.legend()
plt.grid(True)
plt.show()
```



Best max_depth based on Test AUC: 1

Classification Report for Training Set:

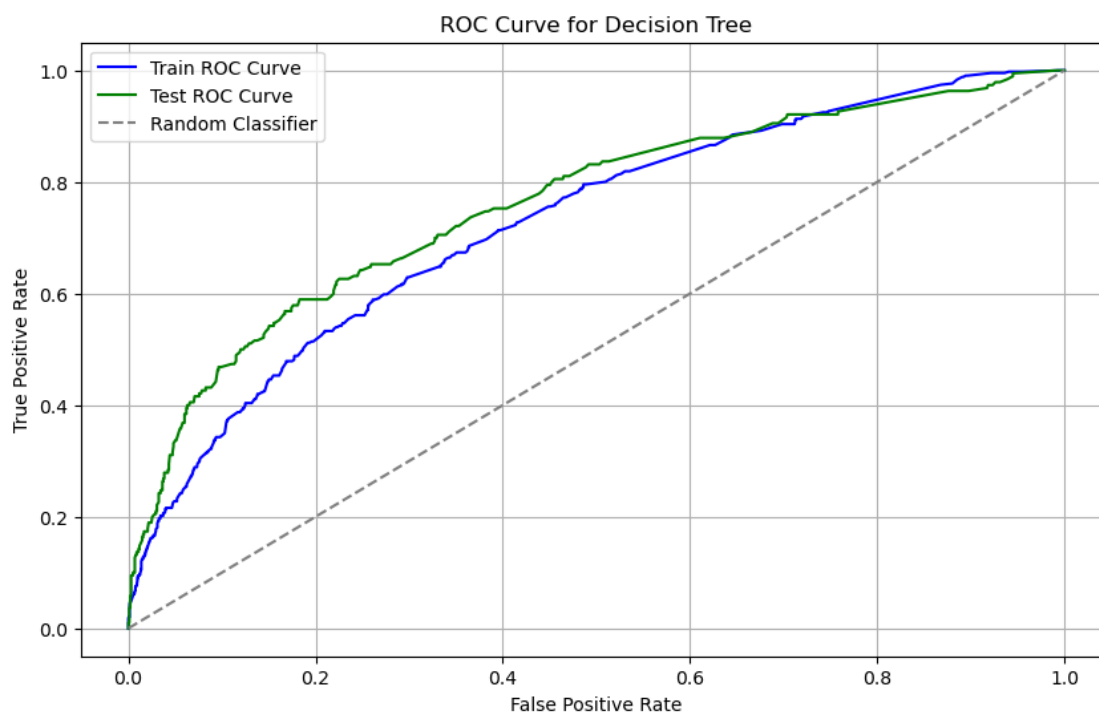
	precision	recall	f1-score	support
0	0.56	0.91	0.69	2390
1	0.75	0.27	0.40	2390
accuracy			0.59	4780
macro avg	0.66	0.59	0.54	4780
weighted avg	0.66	0.59	0.54	4780

Classification Report for Test Set:

	precision	recall	f1-score	support
0	0.88	0.92	0.90	1017
1	0.43	0.33	0.37	190
accuracy			0.83	1207
macro avg	0.66	0.62	0.64	1207
weighted avg	0.81	0.83	0.82	1207

Train AUC: 0.59

Test AUC: 0.62



Graph Interpretation:

Shape of the curve: A ROC curve ideally should be as close to the top-left corner as possible. This indicates that the model has high TPR and low FPR, which is desirable.

Comparison of curves: In this graph, the train ROC curve is slightly above the test ROC curve, suggesting that the model might be slightly overfitting the training data. However, both curves are significantly above the random classifier line, indicating that the model is performing better than random guessing.

Area Under the Curve (AUC): The AUC is a metric that quantifies the overall performance of a model. A higher AUC indicates better performance. In this case, the AUC for both train and test curves would likely be higher than 0.5, indicating that the model is performing better than random guessing.

Overall, the ROC curve suggests that the decision tree model is performing reasonably well on both training and testing data, and is significantly better than a random classifier. However, there might be some overfitting, as indicated by the slight difference between the train and test curves. To improve the model's performance, techniques like regularization or hyperparameter tuning could be considered.

Here are some key take aways from the results above: 1. The model's simplicity: The ideal `max_depth` will be 1 suggesting the model is very simple, reflected in the low performance of the metrics especially for class 1 2. The low recall and F1-scores for class 1 in both the training and test sets suggest that the model struggles to detect the minority class, which is a common issue in imbalanced datasets. 3. Although there isn't a large gap between Train and Test AUC, the overall performance is poor, suggesting that a deeper or more complex model may be necessary to capture the patterns in the data, though care must be taken to avoid overfitting

From the data above it is clear that when we use that data that has been oversampled that all the test performance does not increase across all samples hence there is no positive effect on the test performance.

It is safe to make the assumption that whether we include a different approach of tuning the decision tree the results will be fairly the same in that there will be no change in the test performance of the model.

4.3 LOGISTIC REGRESSION vs DECISION TREE

4.3.1 Analysis of the two models

Choosing the better model between the logistic regression and decision tree (with SMOTE) depends on your specific goals and the trade-offs you're willing to make. Here's a comparison based on the provided results:

4.3.2 Logistic Regression

- **Training Accuracy:** 85.33%
- **Test Accuracy:** 85.50%
- **Precision and Recall:**

- **Class 0:** High precision (0.86), recall (0.99), and f1-score (0.92) on both training and test sets.
- **Class 1:** Lower precision (0.59 train, 0.70 test) and very low recall (0.10 train, 0.14 test).

4.3.3 Decision Tree tuned (max_depth) after SMOTE

- **Training Accuracy:** 59% (This seems lower compared to the logistic regression, possibly due to overfitting or the model complexity.)
- **Test Accuracy:** 83%
- **Precision and Recall:**
 - **Class 0:** Precision (0.88 test), recall (0.92 test), and f1-score (0.90 test) are high.
 - **Class 1:** Lower precision (0.43 test), recall (0.33 test), and f1-score (0.37 test).

4.3.4 Comparative Analysis

1. **Overall Accuracy:** Logistic regression has higher accuracy on both the training and test sets compared to the tuned decision tree.
2. **Class 0 Performance:** Both models perform well on Class 0, but logistic regression maintains consistency across both sets.
3. **Class 1 Performance:**
 - The tuned decision tree shows some improvement in recall and precision for Class 1 compared to the original logistic regression, but the metrics are still relatively low.
 - Logistic regression's performance on Class 1 is quite poor, but the decision tree performs better in recall, suggesting it may better identify some of the minority class samples.

4.3.5 Suitability

- **Logistic Regression:**
 - **Pros:** Consistently high performance on Class 0, high overall accuracy.
 - **Cons:** Struggles with Class 1, leading to a potential under representation of the minority class.
- **Decision Tree with SMOTE:**
 - **Pros:** Improved handling of the minority class (Class 1) due to SMOTE.
 - **Cons:** Lower overall accuracy and training accuracy, potentially overfitting or too simplistic due to max_depth of 1.

4.3.6 Recommendation

If the primary concern is overall accuracy and robust performance on the majority class (Class 0), **logistic regression** is better. However, if improving performance on the minority class (Class 1) is critical and there is a possibility of experimenting with more complex models, the **decision tree with SMOTE** might offer better potential with further tuning.

Trying more advanced models like Random Forests, Gradient Boosting, or Support Vector Machines should be considered, which can potentially offer a better balance between precision and recall for both classes.

To explore further adjustments or other models, the choice depends on the specific application and which metric (precision, recall, F1-score) is more important for the hospital.